


Input layer

Hidden layer 1

Hidden layer 2

Output layer



$$\Phi_{\theta}(x) = W_3 \rho(W_2 \rho(W_1 x + b_1) + b_2) + b_3, \quad \text{trainable parameters } \theta = (W_1, b_1, W_2, b_2, W_3, b_3)$$

ρ activation function, applied component-wise; bias neuron, contributing b_{ℓ} to the following layer